

Experiences in porting Lattice QCD kernels to GPUs

Giannis Koutsou
for

The CyQuarks Group

A. Abdel-Rehim, K. Hadjiyiannakou, A. Vaquero, J. Volmer

Mentors: Matt Norman and Hoshino Tetsuya

Computational-based Science and Technology Research Center (CaSToRC), The Cyprus Institute

Eurohack 2015, July 2015, Lugano, Switzerland




THE CYPRUS
INSTITUTE

CaSToRC

Lattice QCD

QCD: Fundamental theory of strong interactions

- Part of the Standard Model of particle physics
- Responsible for binding quark into protons and neutrons
- Accounts for baryonic mass observed in the universe



■ Post-diction of experimentally well known quantities

- Masses of low-lying hadrons
- Nucleon axial-charge g_A
- Nucleon momentum fraction $\langle x \rangle$
- Nucleon electromagnetic form-factors


“Benchmark” or “Gold-plated” quantities

■ Pre-diction of experimentally less (or not) known quantities

- Nucleon scalar and tensor charges g_S, g_T
- Nucleon sigma terms
- Axial charges of hyperons
- Neutron electric dipole moment
- Precision measurements of proton radius, anomalous magnetic moment of muon (hadronic)


QCD on the lattice

4D space - time lattice



- 4D grid, spacing a , extent L
- Quark fields $\psi(x)$, $\bar{\psi}(x)$, gluon fields $U_\mu(x)$
- Finite $a \rightarrow$ UV cut-off
- Finite $L \rightarrow$ quantized momenta $\frac{2\pi}{L} \vec{n}$


- Monte Carlo simulation for ensembles of gluon fields, with probability $e^{S[U, M^{-1}]}$ (Euclidean time)
- Observables: $\langle O \rangle = \sum_{\{U\}} O(M^{-1}, U_\mu)$, with M^{-1} the quark propagator
- M , discrete covariant derivative (Dirac operator) \rightarrow 4D stencil



Simulation parameters

- Lattice spacing a : take limit $a \rightarrow 0$
- Quark mass m_q : take limit $m_q \rightarrow m_q^{\text{phys}}$
- Volume L : take limit $L \rightarrow \infty$


Simulation cost



Empirical cost formula:

$$C^{\text{sim}} \propto \left(\frac{0.3\text{GeV}}{m_\pi}\right)^{C_m} \left(\frac{L}{2\text{fm}}\right)^{C_L} \left(\frac{0.1\text{fm}}{a}\right)^{C_a}$$

Analysis cost



Typical workflow

■ Pre-process

- Prepare linear system right-hand-sides
- Custom code depending on problem/observable desired
- more maintainable and flexible code required: candidate for OpenACC


■ Production

- Invert for right-hand-sides
- Conjugate-Gradient or variant employed
- Various preconditioning algorithms on the market
- Eigenvalue deflation typically offer $\sim 10\times$ speed-ups
- Use of optimized libraries in CUDA - QUDA package

■ Post-process/analysis

- Generally site-local operations
- Multiplications between arrays of small matrices e.g. arrays of 12×12 complex matrices
- more maintainable and flexible code: candidate for OpenACC

Matrix inversion on multiple right-hand-sides



▪ Eigenvalue deflation

- $\mathcal{O}(1000)$ eigenvectors calculated for $10^7 \times 10^7$ matrix
- CPU code employed up to now


▪ Inversion

- Use of eigenvalues for preconditioning the inversion
- QUDA used, eigenvectors read from disk

▪ Goals during Eurohack: Perform deflation on GPUs

- Main challenge: limited GPU memory
- Solution: Keep eigenvectors on host memory, stream in to GPU for matrix application

Timings



- Strong scaling of smaller-than-production problem size
- Scaling important: more GPUs means more memory meaning more eigenvectors

Analysis kernels

- **Many unrolled sparse-matrix times dense-matrix multiplications**
 - Write as loops to allow vectorisation
 - Generate index arrays of non-zero elements
- **Short, inline and static functions for matrix multiplications**
 - Need to rewrite as macros inline functions
- **Arrays of structs of arrays**
 - Care when transferring C-structures to GPU
- **Mainly using PGI compiler**
 - craycc issue when allocating device pointer

Analysis kernels

```
/* multiply prop by |gamma_0 from the left */
static void
prop_g0_G(_Complex double out[NS*NC][NS*NC], _Complex double in[NS*NC][NS*NC])
{
    out[0][0] = -in[6][0];
    out[0][3] = -in[6][3];
    out[0][6] = -in[6][6];
    out[0][9] = -in[6][9];
    out[3][0] = -in[9][0];
    out[3][3] = -in[9][3];
    out[3][6] = -in[9][6];
    out[3][9] = -in[9][9];
    out[6][0] = -in[0][0];
    out[6][3] = -in[0][3];
    out[9][0] = -in[3][0];
    out[9][3] = -in[3][3];
}

/* multiply prop by |gamma_0 from the left */
static inline void
prop_g0_G(_Complex double out[NS*NC][NS*NC], _Complex double in[NS*NC][NS*NC])
{
    for(int i=0; i<NC*NS; i++)
        for(int j=0; j<NC*NS; j++) {
            out[i][j] = 0.;
            for(int k=0; k<prop_g0_G_nva; k++)
                out[i][j] += (prop_g0_G_val[i][j][k])*in[prop_g0_G_idx[i][j][k][0]][prop_g0_G_idx[i][j][k][1]];
        }
    return;
}

/* multiply prop by *H(|gamma_0) from the right */

/* multiply prop by |gamma_0 from the left */
#define prop_g0_G(out, in)
for(int i=0; i<NC*NS; i++)
    for(int j=0; j<NC*NS; j++) {
        out[i][j] = 0.;
        for(int k=0; k<prop_g0_G_nva; k++)
            out[i][j] += (prop_g0_G_val[i][j][k])*in[prop_g0_G_idx[i][j][k][0]][prop_g0_G_idx[i][j][k][1]];
    }

/* multiply prop by |bar(|gamma_0) from the right */
#define prop_g0_G(out, in)
for(int i=0; i<NC*NS; i++)
    for(int j=0; j<NC*NS; j++) {
```

Analysis kernels

```
#endif
#ifdef QHG_ACC
#pragma acc enter data pcopyin(fwd[0:NC*NS], bwd[0:NC*NS]) async(1)
    for(int cs=0; cs<NC*NS; cs++) {
#pragma acc enter data pcreate(fwd[cs].field[0:lvol_bvol*NC*NS]) async(1)
#pragma acc update device(fwd[cs].field[0:lvol_bvol*NC*NS]) async(1)

#pragma acc enter data pcreate(bwd[cs].field[0:lvol_bvol*NC*NS]) async(1)
#pragma acc update device(bwd[cs].field[0:lvol_bvol*NC*NS]) async(1)
    }


#pragma acc enter data pcreate(corrC[0:lvol*SITE_SIZE]) async(1)

#pragma acc parallel loop gang vector_length(144)          \
    private(F, B, T, gF, fwd_tbc, bwd_tbc, tsrc)          \
    present(fwd, bwd, corrC) async(1)
#endif
    for(int v=0; v<lvol; v++) {
#pragma acc cache(
        F[0:(NC*NS)][0:(NC*NS)], \
        B[0:(NC*NS)][0:(NC*NS)], \
        gF[0:(NC*NS)][0:(NC*NS)], \
        T[0:(NC*NS)][0:(NC*NS)])

#pragma acc loop vector collapse(2)
        prop_load_(F, fwd, v);
#pragma acc loop vector collapse(2)
        prop_load_(B, bwd, v);

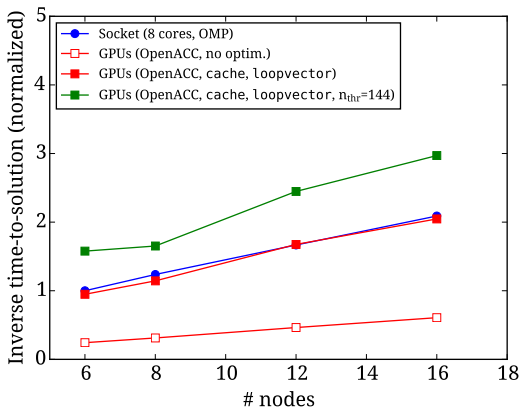
        int t = v/lv3;
        int gt = t + t0;
        if(gt < tsrc) {
#pragma acc loop vector collapse(2)
            prop_scale_(fwd_tbc, F);
#pragma acc loop vector collapse(2)
            prop_scale_(bwd_tbc, B);
        }
    }
}
```

Timings




- Including any explicit `vector_length()` leads to wrong results
- Some speed-up may be possible, e.g. `vector_length(64)` produced a faster kernel

Timings



- Including explicit `vector_length()` leads to wrong results
- Some speed-up may be possible, e.g. various choices of `vector_length()` produced a faster kernel, but results are not always identical

Timings



- ~25% time in memcpy
- Production kernels will be larger \Rightarrow potential for speed-up
- Buggy reduction over threads

Impressions

▪ OpenACC impressions

- Changes to code required (some unpleasant)
 - ▶ Un-unroll manually unrolled loops
 - ▶ Introduction of index arrays
 - ▶ Change inline functions to macros
 - No impact on CPU code
- Maintainability
 - ▶ Can switch between OpenACC/OpenMP with `ifdefs`

▪ Hackathon impressions

- Excellent opportunity to concentrate on code development
- Much more useful than schools which typically include only short example codes
- Mentors at table were critical to go forward
 - . Perhaps an extra day to ease-in would be helpful
 - . Or as suggested to have a communication to start some OpenACC development a couple of weeks before Hackathon
 - . Open to general GPU development (CUDA?)

Thank you!

A. Abdel-Rehim, C. Alexandrou, K. Hadjiyiannakou, A. Vaquero, J. Volmer



EUROPEAN UNION



REPUBLIC OF CYPRUS



Research
Promotion
Foundation



STRUCTURAL FUNDS
of the European Union in Cyprus
our ideas, actions for development

The Project GPU Clusterware (ΤΠΕ/ΠΛΗΡΟ/0311(BIE)/09)