

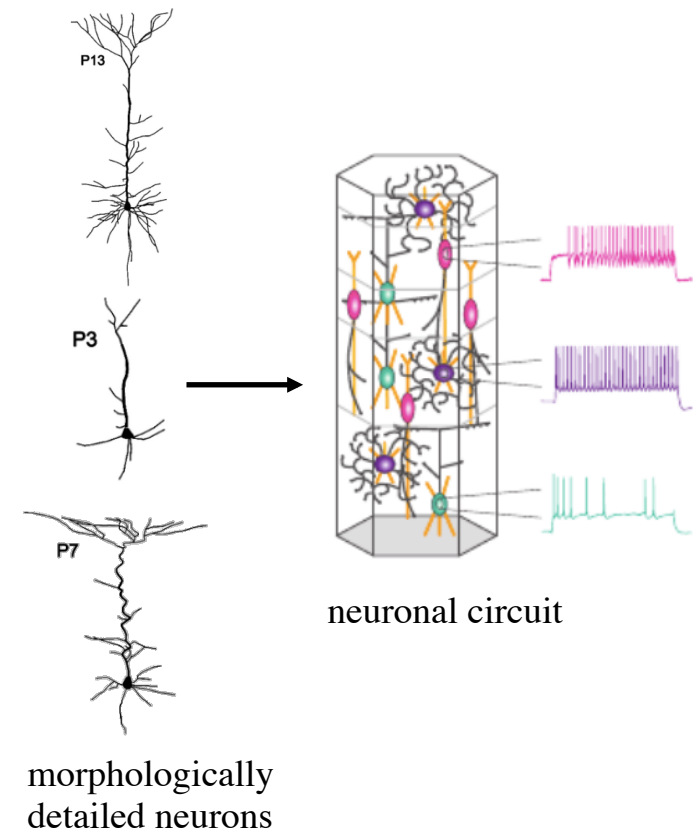
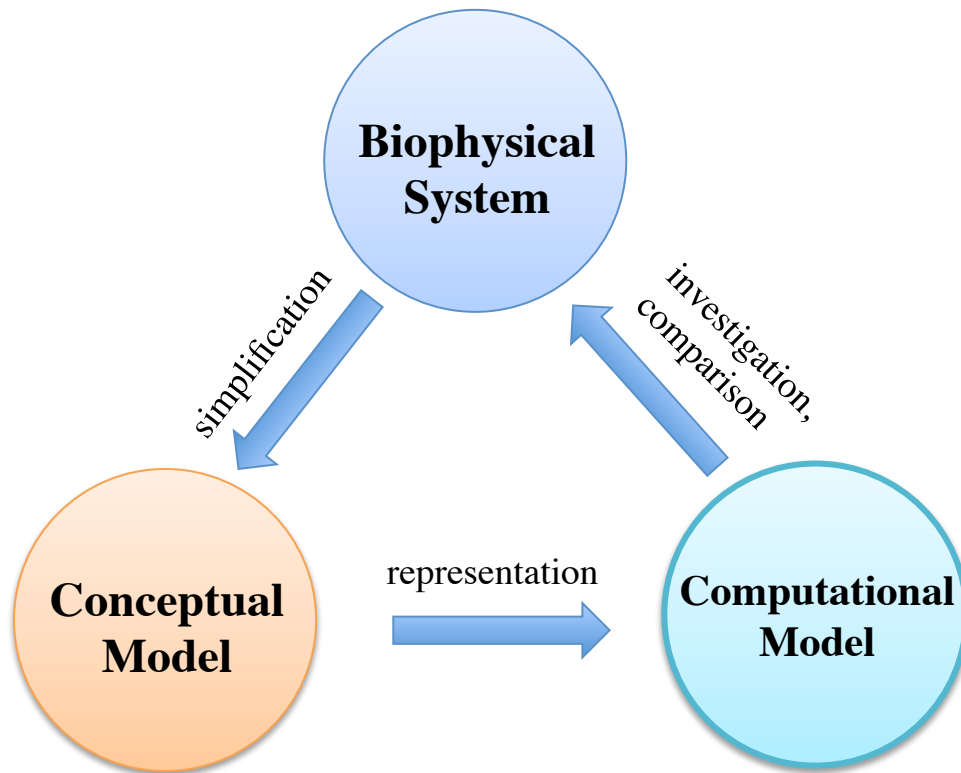
CoreNeuron

Journey to Eurohack 2015

Blue Brain Project team, Ben, Jakob

Comprehensive approach to systematically create unifying models of brain circuits by

- reverse engineering biological components
- construction of math models of the biophysics



Preparation



48 Racks, MIRA, ANL
3.14m threads, 260m neurons
May 2015



28 Racks, JUQUEEN, Juelich
1.8m threads, 160m neurons
Feb 2015



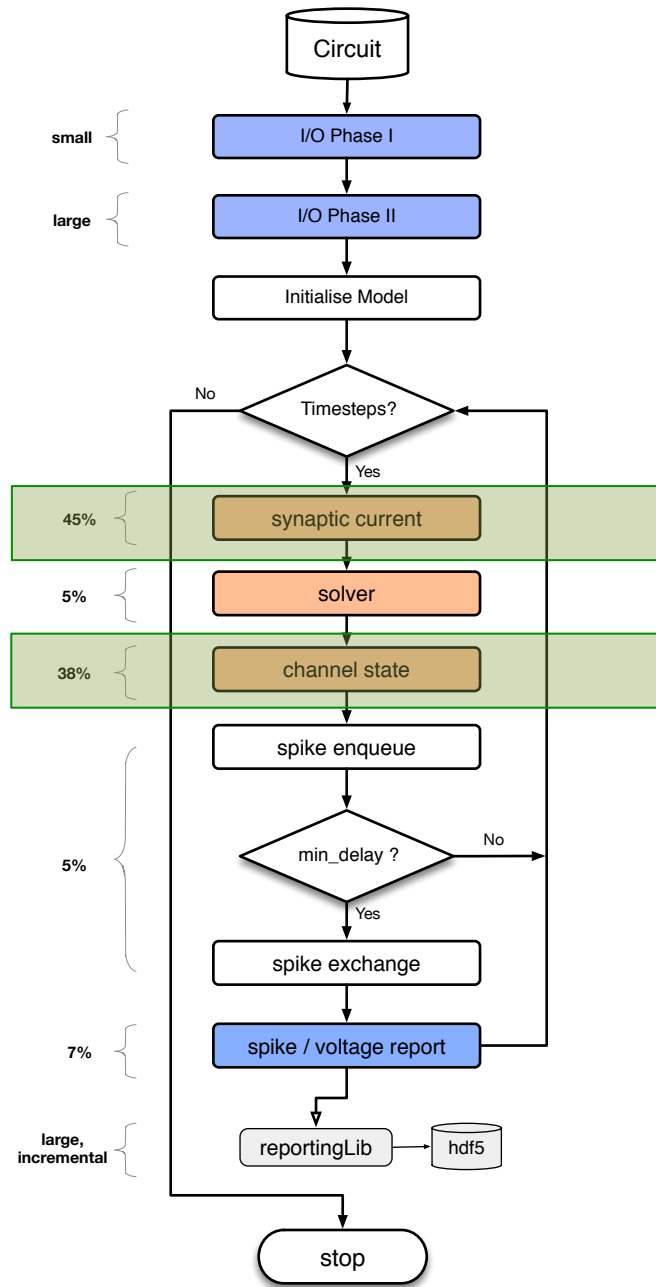
4 Racks, BBP IV, Lugano
262k threads, 10m neurons
Dec 2014

Preparation : Before the workshop

- Preparing simulation dataset
- Compilation with Cray, Intel and PGI
 - fixing Random123 issue with cray
- Profiling on host / sandybridge
 - verify the cpu performance with

What to offload?

CoreNeuron Simulator



- **Stimulus, Solver: 5-8%**
- **Channel State update: 90%**

Porting Challenges

Porting Challenges: Data Structure

```
typedef struct NrnThread {
    double _t;
    double _dt;
    double cj;
    NrnThreadMembList* tml;
    int ncell; /* analogous to old rootnodecount */
    int end; /* 1 + position of last in v_node a
    int id; /* this is nrn_threads[id] */
    int _stop_stepping; /* delivered an all threa

    double* _actual_rhs;
    double* _actual_d;
    double* _actual_a;
    double* _actual_b;
    double* _actual_v;
    double* _actual_area;
    int* _v_parent_index;
    Node** _v_node;
    Node** _v_parent;
    char* _sp13mat; /* handle to general sparse ma
    Memb_list* _ecell_memb_list; /* normally n
    void* _vcv; /* replaces old ccode_instance and r

#ifdef 1
    double _ctime; /* computation time in seconds (
#endif


    NrnThreadBAList* tbl[BEFORE_AFTER_SIZE]
    hoc_List* roots; /* ncell of these */
    Object* userpart; /* the SectionList if this

} NrnThread;
```

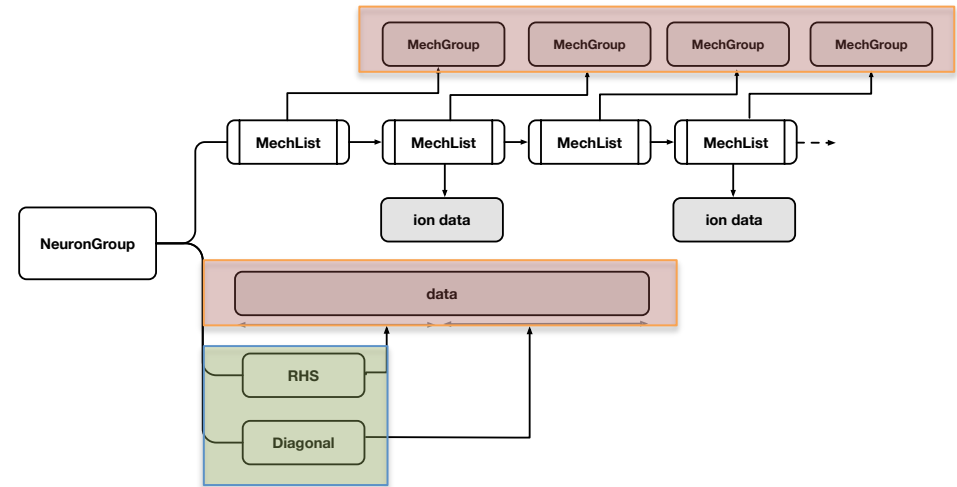
```
typedef struct Memb_list {
    Node** nodelist;
#ifdef CACHEVEC != 0
    /* nodeindices contains all nodes this extension is responsible for,
    * ordered according to the matrix. This allows to access the matrix
    * directly via the nrn_actual_* arrays instead of accessing it in the
    * order of insertion and via the node-structure, making it more
    * cache-efficient */
    int *nodeindices;
#endif /* CACHEVEC */
    double** data;
    Datum** pdata;
    Prop** prop; //DIMENSION IS nodecount
    Datum* _thread; /* thread specific data (when static is no good) */
    int nodecount;
} Memb_list;

typedef union Datum { /* interpreter stack type */
    double val;
    Symbol *sym;
    int i;
    double *pval; /* first used with Eion in NEURON */
    HocStruct Object **pobj;
    HocStruct Object *obj; /* sections keep this to construct a name */
    char **pstr;
    HocStruct hoc_Item* itm;
    hoc_List* lst;
    void* _pvoid; /* not used on stack, see nrnoc/point.c */
} Datum;
```


Challenges: Data Structure & Lots of Kernels



Biologist view:
compartment model



Memory View: In memory
representation of Neurons

```
DERIVATIVE states {
  LOCAL mAlpha, mBeta, mInf, mTau, lv, qt

  qt = 2.952882641412121
  lv = v

  if(lv == -32){
    lv = lv+0.0001
  }

  mAlpha = mAlphaf(lv)
  mBeta = mBetaf(lv)
  mInf = mAlpha/(mAlpha+mBeta)
  mTau = (1/(mAlpha+mBeta))/qt
  m' = (mInf-m)/mTau

  v = lv
}
```


mod2c

.C


Hackathon Development

Two teams

MiniApp Benchmark



Data Struct Workflow



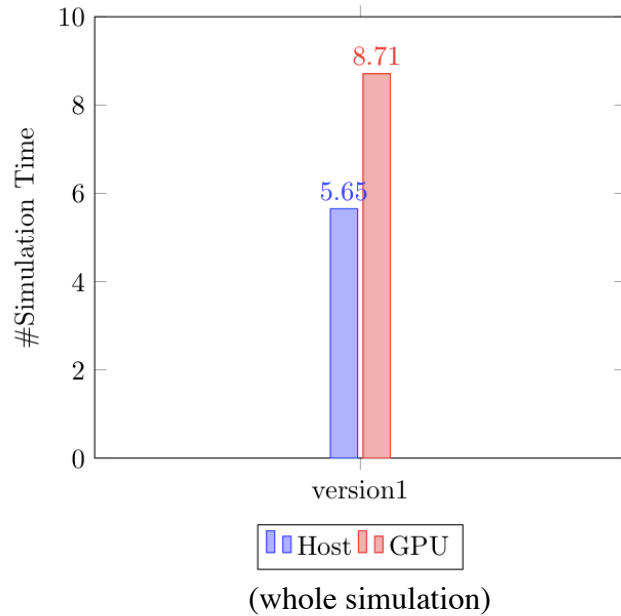
```
_PRAGMA_FOR_VECTOR_LOOP_  
for( i = 0; i < count; i++) {  
  
    int idx = node_index[i];  
    v = vec[idx];  
  
    p3[i] = data[ion_index[i]];  
  
    double gNaTs2 = p0[i]*p1[i]*p1[i]*p1[i]*p2[i];  
    double ina = gNaTs2*(v-p3[i]);  
  
    data[ion_index1[i]] += gNaTs2;  
    data[ion_index2[i]] += ina;  
  
    vec_rhs[idx] -= ina;  
}
```

wrap OpenACC and auto-
vectorisation related pragmas

OpenACC API's to copy the
complex data structure


```
1  
2  for (tml = nt->tml; tml; tml = tml->next) {  
3  
4      /*copy all double data for thread */  
5      d__data = (double *) acc_copyin(nt->_data,  
6                                     nt->_ndata*sizeof(double));  
7  
8      /*update d_nt._data to point to device copy */  
9      acc_memcpy_to_device(&(d_nt->_data), &d__data,  
10                          sizeof(double*));  
11  }
```

Results

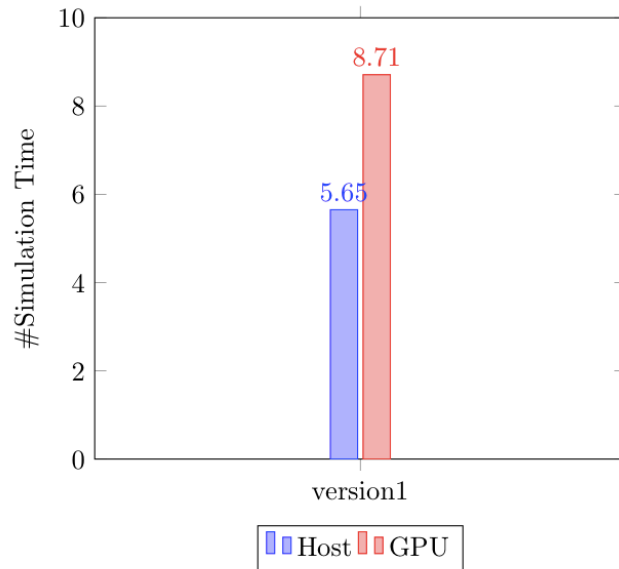


- Single thread offloading kernels from CPU
- Host node 1.5x faster compare to single GPU

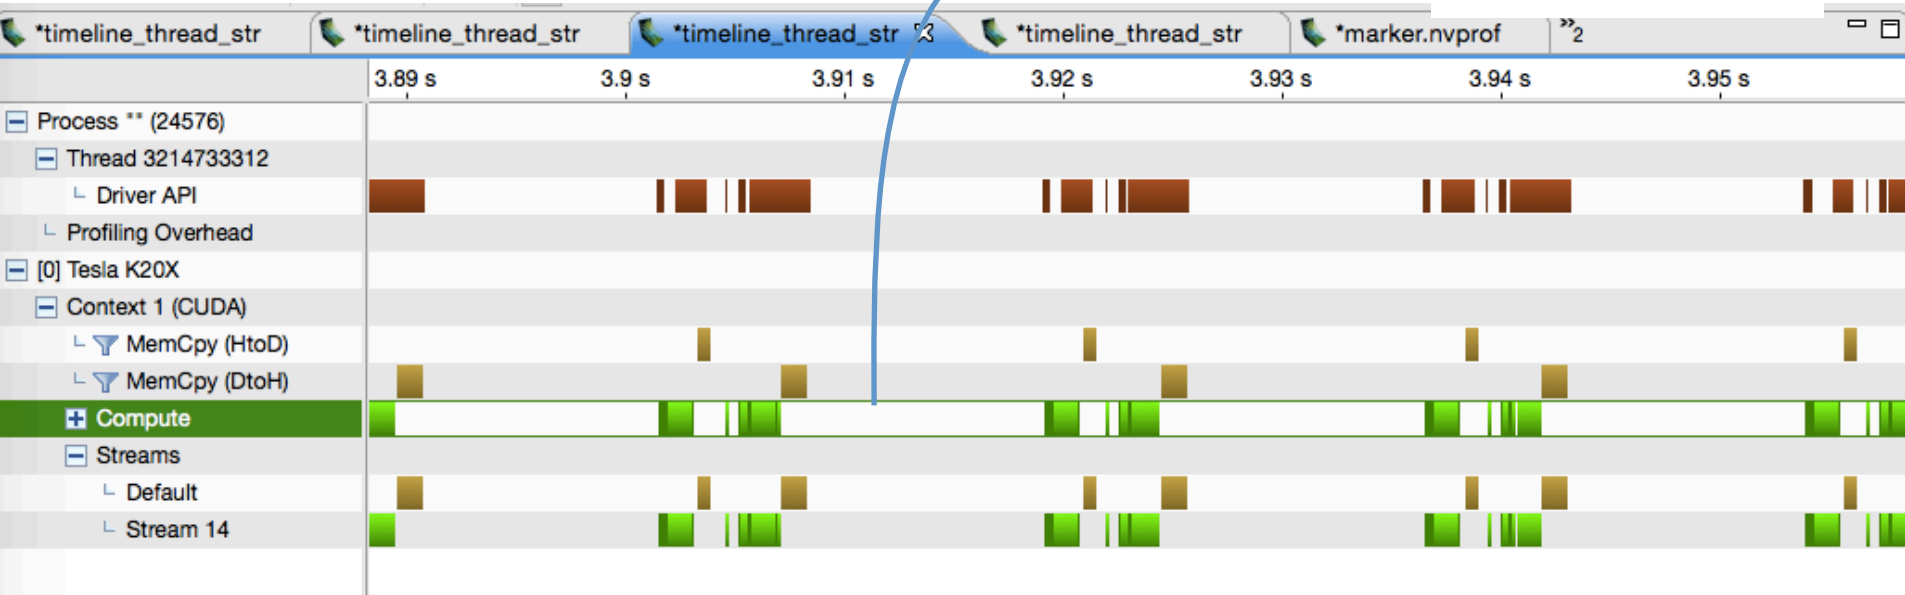
Individual kernels



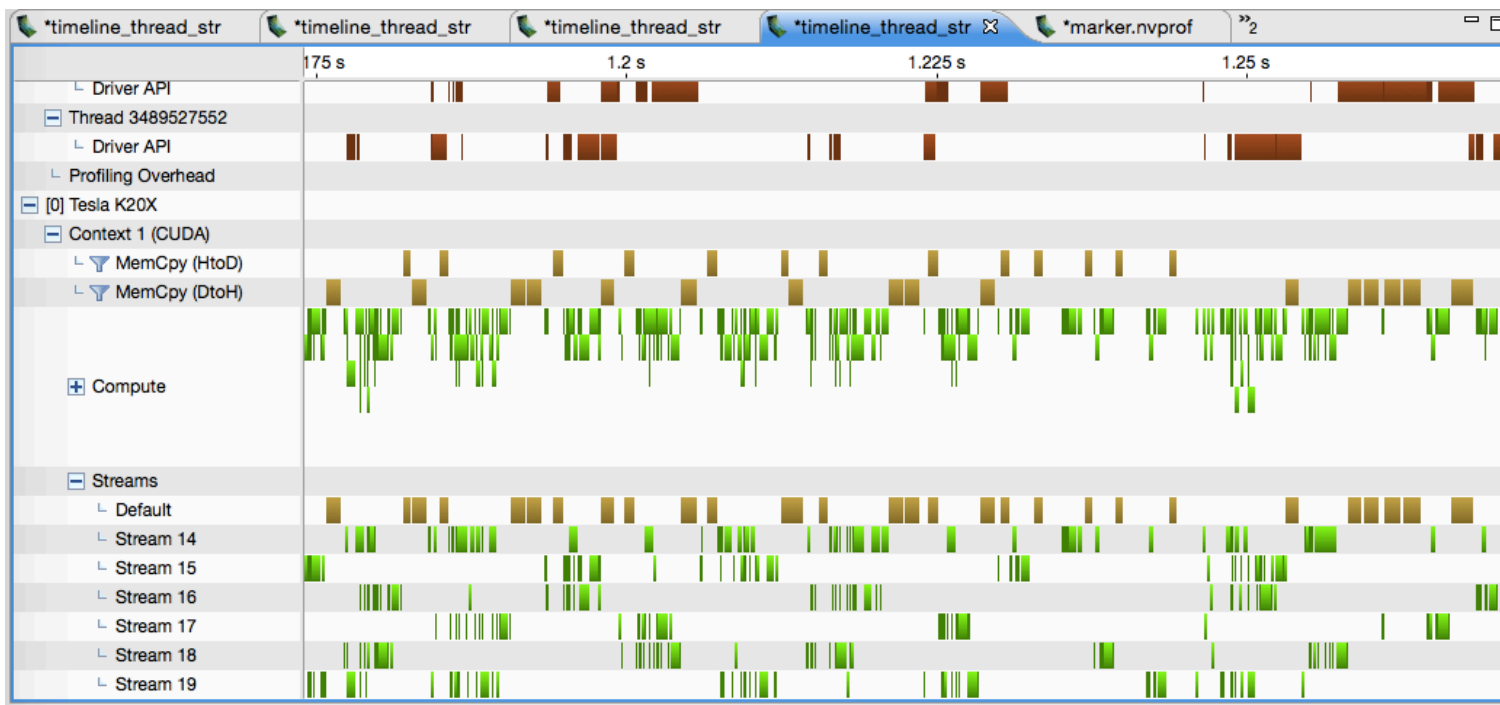
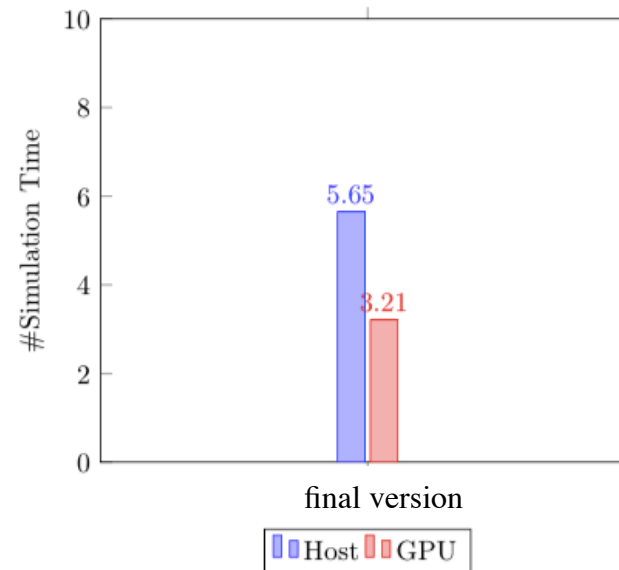
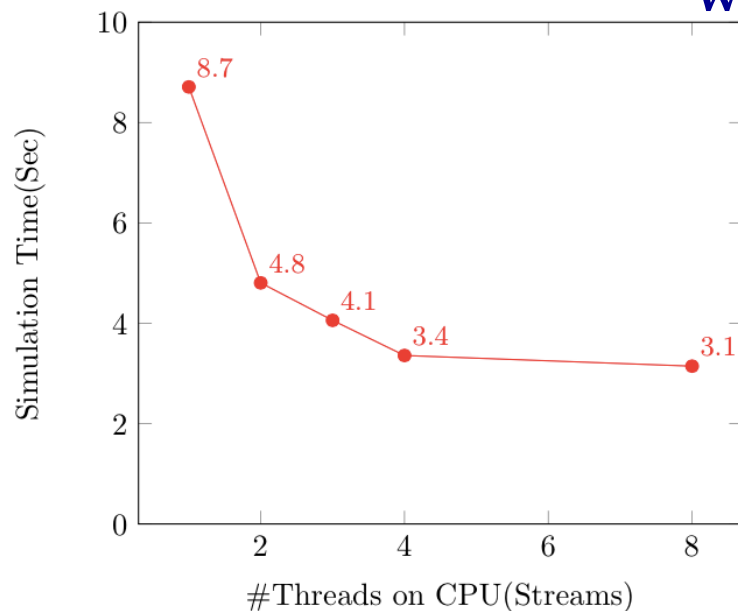
Why GPU code slower?



- Single thread offloading kernels from CPU
- Host node 1.5x faster compare to single GPU



Multiple threads offloading to GPU with streams



Compiler Issues

- Cray
 - `copy(vec[0:asdfgh])`
- PGI
 - vectorization & inlining
- Streams
 - only default stream being used for memcpy?

Next Steps

Ported kernels with OpenACC shows very good speedup compare to CPU (node to node comparison). In order to improve performance:

- Stimulus injection on GPU
- Solver
- Other small routines for spike activity