

# Getting Ready to EuroHack

Alistair Hart (Cray)  
ahart@cray.com

# Getting Ready to EuroHack

- EuroHack starts on Monday July 6th
- One intensive week is a great opportunity
- But you will get a lot further if you do some work first
- **These slides describe some things you can do now**
  - to give you a running start on day 1.



# Learn a bit about OpenACC

- **EuroHack is not a training course**
  - We're all going to learn a lot about OpenACC
  - But we assume that you know the basics already
- **Things you can do:**
  - Watch some CSCS training videos, e.g.
    - <https://www.youtube.com/playlist?list=PL1tk5IGm7zvQOLguaplvYUfkBCVx0GMyW>
  - Read the OpenACC standard (it's not as bad as you think)
    - Download from <http://www.openacc.org>



## Machine access

- **We will use the CSCS machine "Piz Daint"**
  - It has GPU-accelerated nodes
  - It has OpenACC-aware compilers
  - It has a full GPU development environment
- **You should already have access to daint**
  - We will use special course accounts for EuroHack
  - These will avoid your jobs having to wait in the main queue
- **Very first steps:**
  - Make sure you can log in with your course account
  - Make sure you can (and know how to) submit a simple job



# Compiler support

- **daint has two OpenACC compilers**
  - The Cray Compilation Environment (CCE)
  - The PGI compiler
  - (gcc will support OpenACC in the future, but not yet)
- **You will need to be able to compile your code with these**
  - At least one compiler
  - Preferably both, for comparisons and debugging
- **Your next step:**
  - Understand how to use CCE and PGI compilers on daint
  - Develop and test build scripts (Makefiles, etc.) for your code
    - using CCE and PGI
  - Contact CSCS if there are any problems getting your code to compile

# Test cases

- **Code development is an iterative process**
  - Repeated cycles of: edit, compile, run
- **You will need to prepare two testcases to run**
  1. A representative problem
    - Exercises all the relevant parts of the code
    - Should spend the right amount of time in the different parts of the code
    - Reduce this as much as you can (whilst still being representative)
      - weak scale to as few nodes as possible (but not too few)
      - reduce the runtime as much as possible (but not too short)
        - reduce the number of timesteps
        - speed up the initialisation
  2. A rapid development testcase
    - Exercises all the relevant parts of the code
    - Needn't spend the right amount of time in the different parts of the code
    - Make this as quick to run as possible (ideally, a couple of minutes or less)

# Correctness checking

- **A successful OpenACC port is done incrementally**
  - It will not be faster at every step
  - But it should be numerically correct at every step
- **Your next step**
  - Make sure your code has correctness checks (or add them)
    - e.g. known values in output file that don't change (for a given testcase)
      - array checksums, solver residuals...
    - or internal checks in the code
    - The "domain experts" in the team can help with this
  - If you have time, write a small validation script
    - Parses the job output and checks for correctness



# Profiling your code with CrayPAT

- **Profiling is important**
  - not just to measure code performance
  - but also to guide the OpenACC porting process
- **CrayPAT is the Cray profiling tool**
  - It works with all the compilers on daint
  - It profiles CPU and GPU parts of your code
- **CrayPAT works with all the compilers on daint**
  - It has some extra features with CCE, designed for OpenACC porting





# Profiling your code

- **First, a sampling profile is useful:**
  - Build the code:
    - "make clean" to remove all the object files
    - "module load perftools"
    - "make" to build executable EXE
  - Instrument the code
    - "pat\_build -f EXE" to create new executable EXE+pat
  - Now execute EXE+pat using usual jobscript
    - A file or directory is created; name ends in ".xf"
  - Create the profile
    - pat\_report <xf file/directory>
    - Save this report in a text file

# Profiling the right part of your code

- **Don't profile the initialisation**
  - just the main computational steps
  - e.g. timestep loop
- **Use the CrayPAT API to do this**
  - Fortran, C, C++
- **See "man pat\_build" for details**
- **CPP macro: CRAYPAT**
  - Useful so code compiles with and without perf tools
  - CCE: automatically defines CRAYPAT when perf tools loaded
  - PGI: need to manually include -DCRAYPAT when perf tools loaded
- **Repeat your profiles using the API**

```
#ifdef CRAYPAT
include 'pat_apif.h'
#endif
! declarations
#ifdef CRAYPAT
call PAT_record(PAT_STATE_OFF,ipat)
#endif
! initialisation code
#ifdef CRAYPAT
call PAT_record(PAT_STATE_ON,ipat)
#endif
! computation code
#ifdef CRAYPAT
call PAT_record(PAT_STATE_OFF,ipat)
#endif
! finalisation code
```



## Getting a calltree

- **A calltree shows which routines call which**
  - Potentially changes for different testcases
  - Provides a map for the OpenACC port
- **The OpenACC port will start at the child routines**
  - Then work back up the calltree
- **Generate the calltree from existing CrayPAT data:**
  - `pat_report -O calltree -T <xf file/directory>`
  - Save this report in a text file
- **-T includes all routines in the output**
  - Not just those taking more than 2% (the default)



## Loop level profiling

- **Profiling is typically at the routine level**
- **Profiling at the loop level is very useful for OpenACC**
  - Which loopnests take the most time?
  - How many iterations do these loopnests typically have?
  - How much work is there per iteration
- **CrayPAT supports loop-level profiling with CCE only**
  - module load perftools
  - compile with new compiler flag "-h profile\_generate"
  - pat\_build and execute the code as before
  - pat\_report now gives loop-level profiling information
  - Save this report in a text file



# Summary

Steps you can take now to get a running start at EuroHack:

- 1. Learn about OpenACC**
  - 2. Port your existing code to CCE and PGI compilers**
  - 3. Develop suitable testcases**
  - 4. Develop suitable correctness checks**
  - 5. Generate profiles of your existing code with CrayPAT**
    - CrayPAT API
    - Flat profile
    - Calltree
    - Loop level profiling with CCE
- 
- **You should definitely aim to complete steps 1-4.**
  - **It's even better if you can do step 5 as well.**